

```
import numpy as np
import pandas as pd
import seaborn as sb
import matplotlib.pyplot as plt
import random as rand
from sklearn.neighbors import KNeighborsRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
```

```
df = pd.read_csv("car_prices.csv")
df.head()
```



	year	make	model	trim	body	transmission	vin	state	condition
0	2015	Kia	Sorento	LX	SUV	automatic	5xyktca69fg566472	ca	
1	2015	Kia	Sorento	LX	SUV	automatic	5xyktca69fg561319	ca	
2	2014	BMW	3 Series	328i SULEV	Sedan	automatic	wba3c1c51ek116351	ca	4
3	2015	Volvo	S60	T5	Sedan	automatic	yv1612tb4f1310987	ca	4
4	2014	BMW	6 Series Gran Coupe	650i	Sedan	automatic	wba6b2c57ed129731	ca	4

```
['e'].str.lower()
['del'].str.lower()
```

```
f['make'] == 'ford') & (df['model'] == 'f150') | (df['model'] == 'f-150')).copy()
```

```
index(drop=True, inplace=True)
```

```
lums=['make', 'model'], inplace=True)
```

```
ad())
```

```

year    trim    body transmission    vin    state    condition \
0  2012    XLT    SuperCrew    NaN    1ftew1cm9ckd05952    ca    46.0
1  2012    FX2    SuperCrew    automatic    1ftfw1ct0cfb64807    ca    46.0
2  2012    XLT    SuperCrew    automatic    1ftfw1et3ckd61619    ca    39.0
3  2012    XLT    SuperCrew    automatic    1ftfw1ef9cfc79834    ca    35.0
4  2012    Lariat    SuperCab    automatic    1ftfx1ef6cfc80260    ca    41.0

odometer    color    interior    seller    mmr \
0    51189.0    gray    gray    ford motor credit company llc pd    19700.0
1    28687.0    gray    black    ford motor credit company llc pd    26700.0
2    27465.0    black    gray    ford motor credit company llc pd    26300.0
3    93858.0    white    gray    aaero sweet company    18900.0
4    46444.0    red    black    ford motor credit company llc pd    26300.0

sellingprice    saledate
0    20200.0    Thu Dec 18 2014 12:00:00 GMT-0800 (PST)
1    30500.0    Thu Dec 18 2014 12:00:00 GMT-0800 (PST)
2    28800.0    Thu Dec 18 2014 12:00:00 GMT-0800 (PST)
3    19500.0    Wed Dec 17 2014 12:15:00 GMT-0800 (PST)
4    25600.0    Thu Dec 18 2014 12:00:00 GMT-0800 (PST)

```

```
def reservoir_sampling(m, data):
```

```
    if m > len(data):
```

```
        m = len(data)
```

```
    rsdata = data.copy()
```

```
    reservoirindex = []
```

```
    for i in range(m):
```

```
        reservoirindex.append(i)
```

```
    for i in range(m, len(data)):
```

```
        j = rand.randint(0, i-1)
```

```
        if (j < m):
```

```
            reservoirindex[j] = i
```

```
    return rsdata.iloc[reservoirindex]
```

```
numerical_df = ford_f150_df.select_dtypes(include=['number'])
```

```
sampled_data = reservoir_sampling(14479, numerical_df)
```

```
cleaned_df = sampled_data.dropna()
```

```
print(cleaned_df.head())
```

```

year    condition    odometer    mmr    sellingprice
0  2012    46.0    51189.0    19700.0    20200.0
1  2012    46.0    28687.0    26700.0    30500.0

```

2	2012	39.0	27465.0	26300.0	28800.0
3	2012	35.0	93858.0	18900.0	19500.0
4	2012	41.0	46444.0	26300.0	25600.0

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsRegressor

numeric_features = ['year', 'condition', 'odometer', 'sellingprice']

for feature in numeric_features:
    cleaned_df.loc[:, feature] = pd.to_numeric(cleaned_df[feature], errors='coerce')
    cleaned_df.loc[:, feature] = cleaned_df[feature].fillna(cleaned_df[feature].mean)

cleaned_df = cleaned_df.replace([np.inf, -np.inf], np.nan)

for feature in numeric_features:
    cleaned_df.loc[:, feature] = cleaned_df[feature].fillna(cleaned_df[feature].mean)

X = cleaned_df[numeric_features].values
y = cleaned_df['mmr'].values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_stat

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

knn = KNeighborsRegressor(n_neighbors=4)
knn.fit(X_train_scaled, y_train)

predicted_values = knn.predict(X_test_scaled)

errors = predicted_values - y_test
errors = abs(errors)

plt.figure(figsize=(10, 6))
plt.scatter(y_test, predicted_values, color='blue')

plt.xlabel('predicted_values')
plt.ylabel('y_test')
plt.title('Scatter Plot of predicted_values vs test values')

plt.xlim(0, 50000)
plt.ylim(0, 50000)

plt.show()
```

